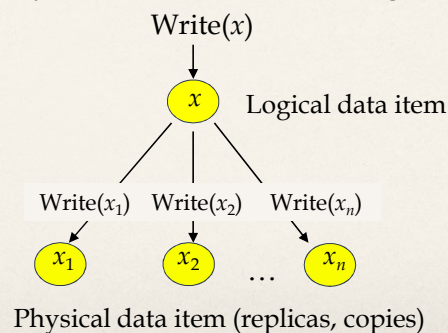


## Data Replication

- Benefits of Replication
  - System Availability
    - ♦ May remove single points of failure by replicating data
  - Performance
    - ♦ Replication enables us to locate the data closer to their access point thereby reducing response point
  - Scalability
    - ♦ Replication allows for a way to support systems growth with acceptable response time
  - Application requirements
    - ♦ Some applications may wish to maintain multiple data copies as part of their operational specifications

## Execution Model

- There are physical copies of logical objects in the system.
- Operations are specified on logical objects, but translated to operate on physical objects.
- One-copy equivalence
  - The effect of transactions performed by clients on replicated objects should be the same as if they had been performed on a single set of objects.



## Issues in Design of Replication Protocols

- 1. Database design:
  - Distributed databases may be fully or partially replicated.
- 2. Database consistency
  - Mutual consistency:
    - ♦ A replicated database is said to be in mutually consistent (or strong consistent) state if all replicas of each of its data items have identical values.
  - Weak consistency
    - ♦ Eventual consistency: the copies are not identical when update transaction completes, but they eventually converge to the same value

## Issues in Design of Replication Protocols

- 3. Where updates are performed
  - Centralized
  - Distributed
- 4. Update propagation techniques – how do we propagate updates to one copy to the other copies?
  - Eager
  - Lazy
- 5. Degree of replication transparency

# Transactional Consistency

- How can we guarantee that the global execution history over replicated data is serializable?
- One-copy serializability (1SR)
  - The effect of transactions performed by clients on replicated objects should be the same as if they had been performed *one at-a-time* on a single set of objects.
- Weaker forms are possible
  - Snapshot isolation
  - RC-serializability

Distributed DBMS

©M. T. Özsu &amp; P. Valduriez

Ch.13/5

## Example 1

<u>Site A</u>	<u>Site B</u>	<u>Site C</u>
$x$	$x, y$	$x, y, z$
$T_1$ : $x \leftarrow 20$ Write( $x$ ) Commit	$T_2$ : Read( $x$ ) $y \leftarrow x+y$ Write( $y$ ) Commit	$T_3$ : Read( $x$ ) Read( $y$ ) $z \leftarrow (x*y)/100$ Write( $z$ ) Commit

Consider the three histories:

$$H_A = \{W_1(x_A), C_1\}$$

$$H_B = \{W_1(x_B), C_1, R_2(x_B), W_2(y_B), C_2\}$$

$$H_C = \{W_2(y_C), C_2, R_3(x_C), R_3(y_C), W_3(z_C), C_3, W_1(x_C), C_1\}$$

Global history non-serializable:  $H_B: T_1 \rightarrow T_2, H_C: T_2 \rightarrow T_3 \rightarrow T_1$

Mutually consistent: Assume  $x_A = x_B = x_C = 10, y_B = y_C = 15, z_C = 7$  to begin; in the end  $x_A = x_B = x_C = 20, y_B = y_C = 35, z_C = 7$

Distributed DBMS

©M. T. Özsu &amp; P. Valduriez

Ch.13/6

## Example 2

Site A	Site B
$x$	$x$
$T_1$ : Read( $x$ ) $x \leftarrow x+5$ Write( $x$ ) Commit	$T_2$ : Read( $x$ ) $x \leftarrow x*10$ Write( $x$ ) Commit

Consider the two histories:

$$H_A = \{R_1(x_A), W_1(x_A), C_1, W_2(x_A), C_2\}$$

$$H_B = \{R_2(x_B), W_2(x_B), C_2, W_1(x_B), C_1\}$$

Global history non-serializable:  $H_A: T_1 \rightarrow T_2, H_B: T_2 \rightarrow T_1$

Mutually inconsistent: Assume  $x_A = x_B = 1$  to begin; in the end  $x_A = 6$ ,  $x_B = 10$

Distributed DBMS

©M. T. Özsu &amp; P. Valduriez

Ch.13/7

## Update Management Strategies

- Depending on when the updates are propagated
  - Eager
  - Lazy
- Depending on where the updates can take place
  - Centralized
  - Distributed

	Centralized	Distributed
Eager		
Lazy		

Distributed DBMS


©M. T. Özsu &amp; P. Valduriez

Ch.13/8



# Eager Replication

- Changes are propagated within the scope of the transaction making the changes.




Distributed DBMS

©M. T. Özsu &amp; P. Valduriez

Ch.13/9

# Lazy Replication

- Lazy replication first executes the updating transaction on one copy. After the transaction commits, the changes are propagated to all other copies (**refresh transactions**)




Distributed DBMS

©M. T. Özsu &amp; P. Valduriez

Ch.13/10

## Centralized

- There is only one copy which can be updated (the **master**), all others (**slave copies**) are updated reflecting the changes to the master.




Distributed DBMS

©M. T. Özsu &amp; P. Valduriez

Ch.13/11

## Distributed

- Changes can be initiated at any of the copies. That is, any of the sites which owns a copy can update the value of the data item.



Distributed DBMS

©M. T. Özsu &amp; P. Valduriez

Ch.13/12

# Forms of Replication

## Eager

- + No inconsistencies (identical copies)
- + Reading the local copy yields the most up to date value
- + Changes are atomic
- A transaction has to update all sites
  - Longer execution time
  - Lower availability

## Lazy

- + A transaction is always local (good response time)
- Data inconsistencies
- A local read does not always return the most up-to-date value
- Changes to all copies are not guaranteed
- Replication is not transparent

## Centralized

- + No inter-site synchronization is necessary (it takes place at the master)
- + There is always one site which has all the updates
- The load at the master can be high
- Reading the local copy may not yield the most up-to-date value

## Distributed

- + Any site can run a transaction
- + Load is evenly distributed
- Copies need to be synchronized