

## The MapReduce Paradigm

- Platform for reliable, scalable parallel computing
- Abstracts issues of distributed and parallel environment from programmer.
- Runs over distributed file systems
  - Google File System (GFS)
  - Hadoop File System (HDFS)

## Problem Scope

- *MapReduce* is a parallel programming model for data processing
- The power of MapReduce lies in its ability to scale to 100s or 1000s of computers, each with several processor cores
- How large an amount of work?
  - Web-Scale data on the order of 100s of GBs to TBs or PBs
  - It is likely that the input data set will not fit on a single computer's hard drive
  - Hence, a distributed file system (e.g., Google File System- GFS) is typically required

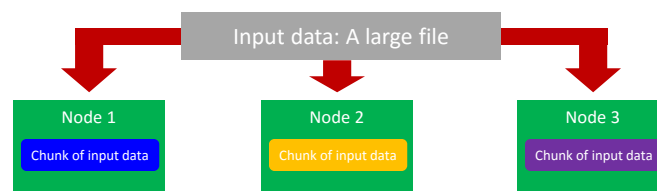
## Commodity Clusters

- MapReduce is designed to efficiently process large volumes of data by connecting many commodity computers together to work in parallel
- MapReduce divides the workload into multiple *independent tasks* and schedule them across cluster nodes
- A work performed by each task is done *in isolation* from one another

3

## Data Distribution

- In a MapReduce cluster, data is distributed to all the nodes of the cluster as it is being loaded in
- An underlying distributed file systems (e.g., GFS) splits large data files into chunks which are managed by different nodes in the cluster



- Even though the file chunks are distributed across several machines, they form *a single namespace*

4

## Functional Abstractions Hide Parallelism

- Map and Reduce
- Functions borrowed from functional programming languages (eg. Lisp)
- Map()
  - Process a key/value pair to generate intermediate key/value pairs•
- Reduce()
  - Merge all intermediate values associated with the same key

11/23/10

## Relational Databases vs. MapReduce

- Relational databases:
  - Multipurpose: analysis and transactions; batch and interactive
  - Data integrity via ACID transactions
  - Lots of tools in software ecosystem (for ingesting, reporting, etc.)
  - Supports SQL (and SQL integration, e.g., JDBC)
  - Automatic SQL query optimization
- MapReduce (Hadoop):
  - Designed for large clusters, fault tolerant
  - Data is accessed in “native format”
  - Supports many query languages
  - Programmers retain control over performance
  - Open source

Source: O'Reilly Blog post by Joseph Hellerstein (11/19/2008)

## MapReduce

- Map: Inputs a key/value pair
  - Key is a reference to the input value
  - Value is the data set on which to operate
- Reduce:
  - Starts with intermediate Key / Value pairs
  - Ends with finalized Key / Value pairs

## MapReduce

- MapReduce programs are executed in two main phases, called mapping and reducing:
  - ▶ Map: the map function is written to convert input elements to key-value pairs.
  - ▶ Reduce: the reduce function is written to take pairs consisting of a key and its list of associated values and combine those values in some way.

## WordCount in MapReduce

- Map:
  - ▶ For a pair  $\langle k1, \text{document} \rangle$  produce a sequence of pairs  $\langle \text{token}, 1 \rangle$ , where token is a token/word found in the document.
- Reduce
  - ▶ For a pair  $\langle \text{word}, \text{list}(1, 1, \dots, 1) \rangle$  sum up all ones appearing in the list and return  $\langle \text{word}, \text{sum} \rangle$ , where sum is the sum of ones.

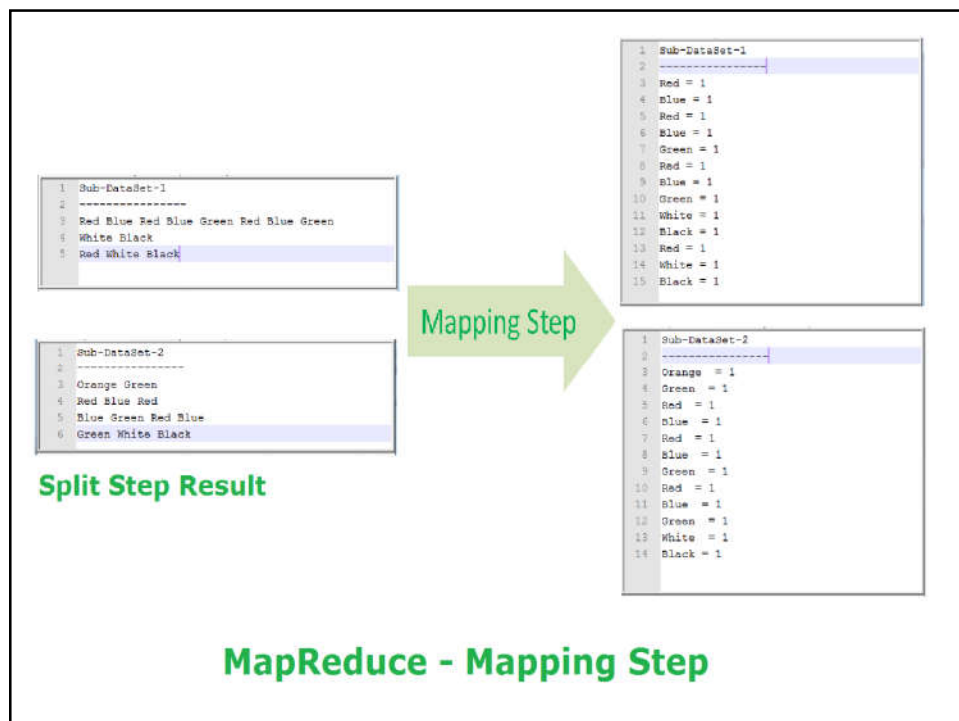
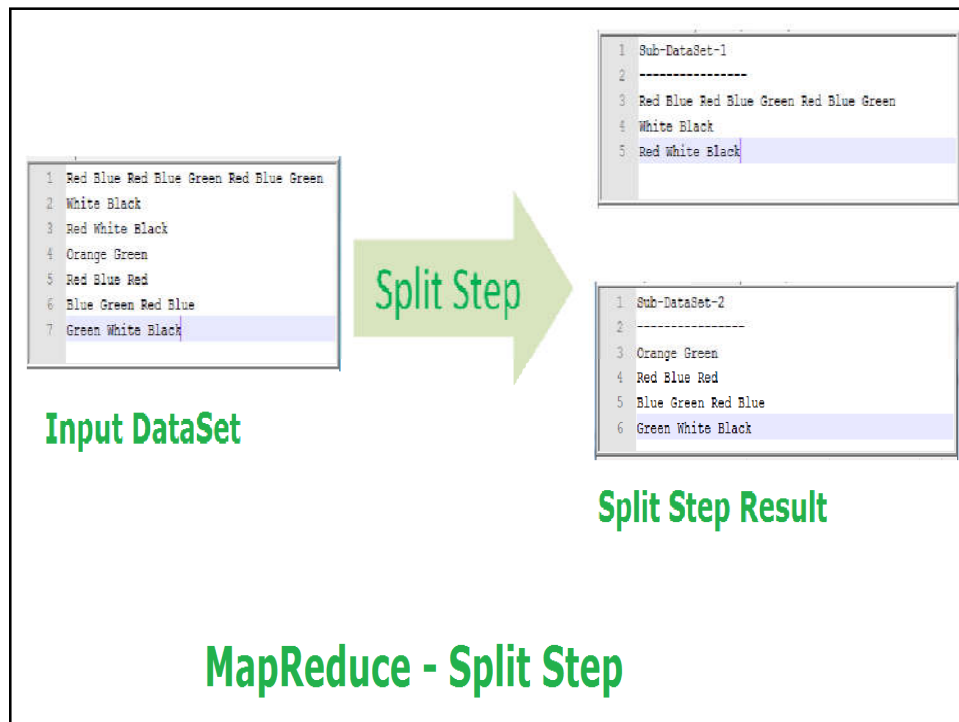
### Example

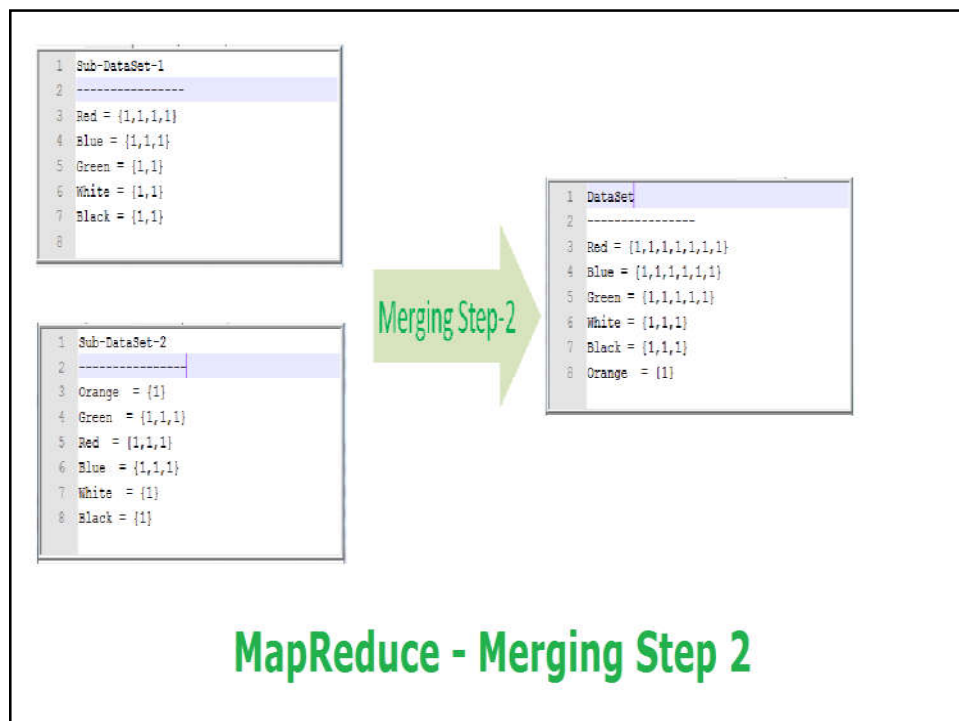
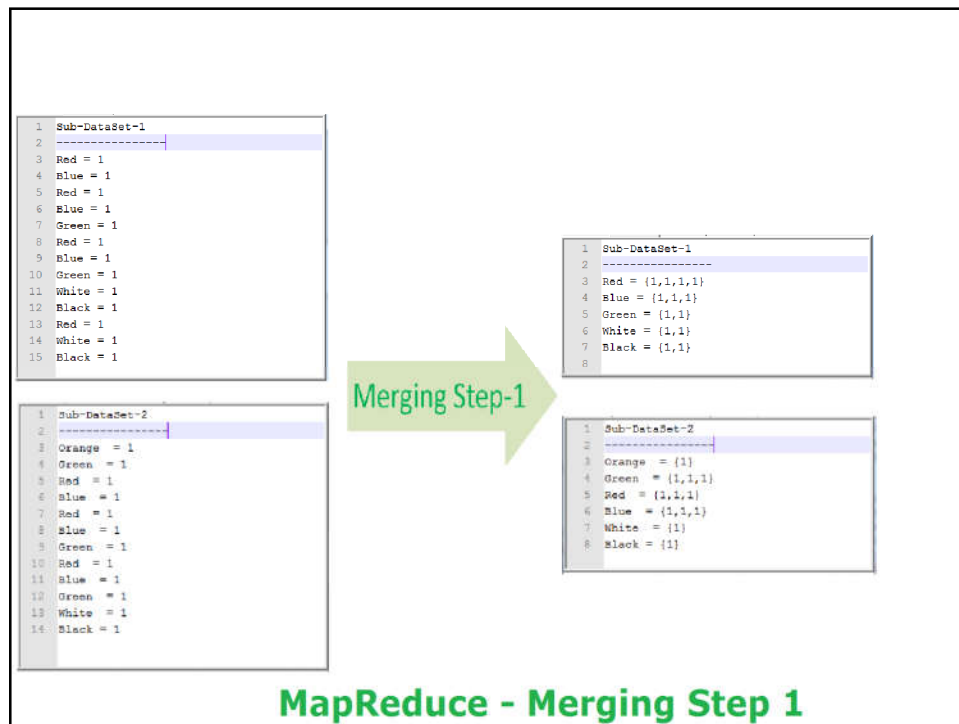
- Count the number of occurrences of each word available in a DataSet given below.

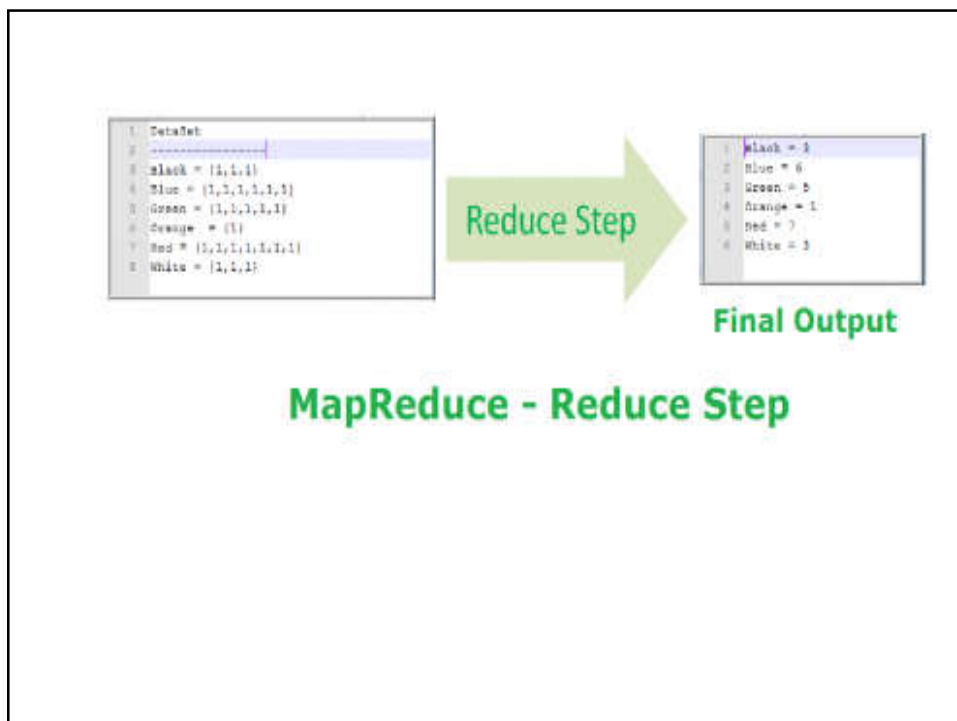
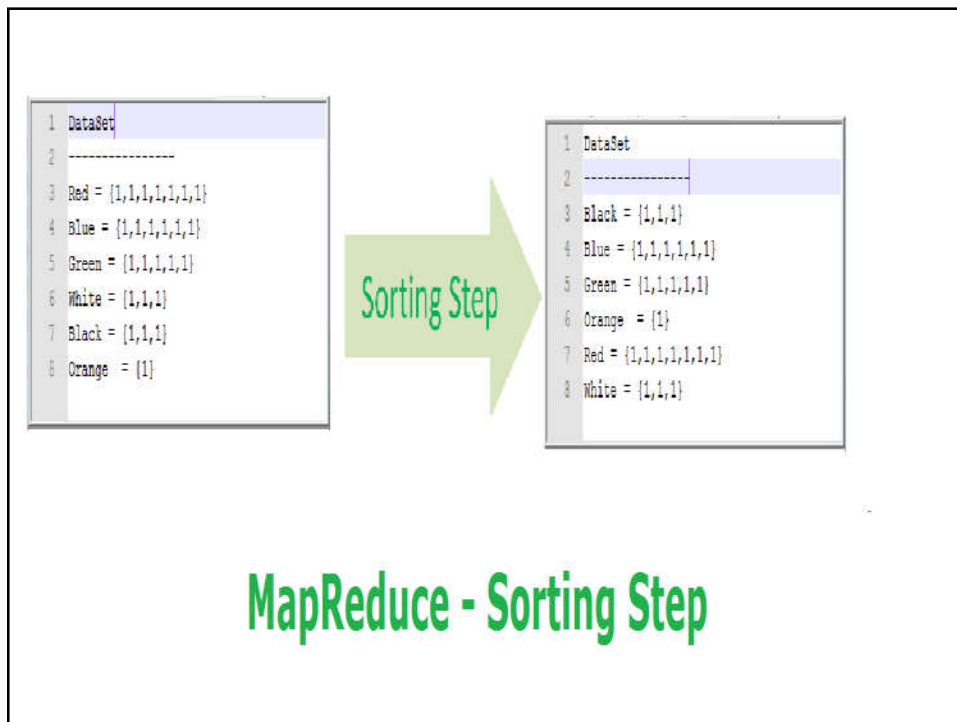
```
1 Red Blue Red Blue Green Red Blue Green
2 White Black
3 Red White Black
4 Orange Green
5 Red Blue Red
6 Blue Green Red Blue
7 Green White Black
```

```
1 Black = 3
2 Blue = 6
3 Green = 5
4 Orange = 1
5 Red = 7
6 White = 3
```

**Final Ouput**









## MapReduce Algorithms for Relational Operations

- Selection
- Projection
- Union, and Intersection
- Natural Join

## MapReduce Algorithms for Relational Operations

- $R, S$  - relation
- $t, t'$  - a tuple
- $C$  - a condition of selection
- $A, B, C$  - subset of attributes
- $a, b, c$  - attribute values for a given subset of attributes

## Selection

Selections really do not need the full power of map-reduce. They can be done most conveniently in the map portion alone, although they could also be done in the reduce portion alone

$$\sigma_C(R)$$

- **Map:** For each tuple  $t$  in  $R$ , test if it satisfies  $C$ . If so, produce the key-value pair  $(t, t)$ . That is, both the key and value are  $t$ .
- **Reduce:** The Reduce function is the identity. It simply passes each key-value pair to the output.

## Selection

Selections really do not need the full power of map-reduce. They can be done most conveniently in the map portion alone, although they could also be done in the reduce portion alone. Here is a map-reduce implementation of selection  $\sigma_C(R)$ .

**The Map Function:** For each tuple  $t$  in  $R$ , test if it satisfies  $C$ . If so, produce the key-value pair  $(t, t)$ . That is, both the key and value are  $t$ .

**The Reduce Function:** The Reduce function is the identity. It simply passes each key-value pair to the output.

Note that the output is not exactly a relation, because it has key-value pairs. However, a relation can be obtained by using only the value components (or only the key components) of the output.

## Projection

Projection is performed similarly to selection, because projection may cause the same tuple to appear several times, the Reduce function must eliminate duplicates. We may compute  $\pi_S(R)$  as follows.

**The Map Function:** For each tuple  $t$  in  $R$ , construct a tuple  $t'$  by eliminating from  $t$  those components whose attributes are not in  $S$ . Output the key-value pair  $(t', t')$ .

**The Reduce Function:** For each key  $t'$  produced by any of the Map tasks, there will be one or more key-value pairs  $(t', t')$ . The Reduce function turns  $(t', [t', t', \dots, t'])$  into  $(t', t')$ , so it produces exactly one pair  $(t', t')$  for this key  $t'$ .

## Union

- **Map:** Turn each input tuple  $t$  either from relation  $R$  or  $S$  into a key-value pair  $(t, t)$ .
- **Reduce:** Associated with each key  $t$  there will be either one or two values. Produce output  $(t, t)$  in either case.

## Intersection

- **Map:** Turn each input tuple  $t$  either from relation  $R$  or  $S$  into a key-value pair  $(t, t)$ .
- **Reduce:** If key  $t$  has value list  $[t, t]$ , then produce  $(t, t)$ . Otherwise, produce nothing.

## Example: Join By Map-Reduce

- **Compute the natural join  $R(A,B) \bowtie S(B,C)$**
- $R$  and  $S$  are each stored in files
- Tuples are pairs  $(a,b)$  or  $(b,c)$

A	B		B	C		A	C
a <sub>1</sub>	b <sub>1</sub>	$\bowtie$	b <sub>2</sub>	c <sub>1</sub>	$=$	a <sub>3</sub>	c <sub>1</sub>
a <sub>2</sub>	b <sub>1</sub>		b <sub>2</sub>	c <sub>2</sub>		a <sub>3</sub>	c <sub>2</sub>
a <sub>3</sub>	b <sub>2</sub>		b <sub>3</sub>	c <sub>3</sub>		a <sub>4</sub>	c <sub>3</sub>
a <sub>4</sub>	b <sub>3</sub>						
<b>R</b>			<b>S</b>				

J. Leskovec, A. Rajaraman, J. Ullman: Mining of Massive Datasets, <http://www.mmms.org>

24

## Join in MapReduce (Reduce-side Join)

---

- Assume to have two relations:  $R(A, B)$  and  $S(B, C)$ 
  - We must find tuples that agree on their  $B$  components
- A MapReduce implementation of Natural Join
  - Map:        For a tuple  $(a,b)$  in  $R$  emit a key/value pair  $(b, ('R',a))$   
              For a tuple  $(b,c)$  in  $S$ , emit a key/value pair  $(b, ('S',c))$
  - Reduce:     If key  $b$  has value list  $[('R',a),('S',c)]$ , emit a key/value pair  $(b, (a,b,c))$