

# MAXIMUM FLOW

- In the maximum-flow problem, we wish to compute the greatest rate at which material can be shipped from the source to the sink without violating any capacity constraints.
- It is one of the simplest problems concerning flow networks.

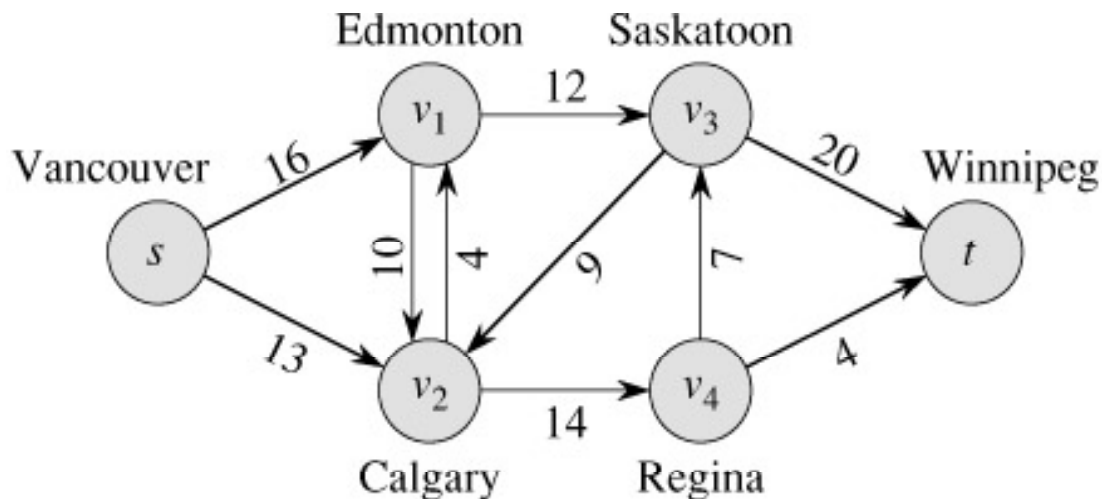
# What is Network Flow ?

- Flow network is a directed graph  $G=(V,E)$  such that each
- edge has a non-negative capacity  $c(u,v) \geq 0$ .
- Two distinguished vertices exist in  $G$  namely :
- Source (denoted by  $s$ ) : In-degree of this vertex is 0.
- Sink (denoted by  $t$ ) : Out-degree of this vertex is 0.

Flow in a network is an integer-valued function  $f$  defined On the edges of  $G$  satisfying  $0 \leq f(u,v) \leq c(u,v)$ , for every Edge  $(u,v)$  in  $E$ .

# What is Network Flow ?

- Each edge  $(u,v)$  has a non-negative capacity  $c(u,v)$ .
  - If  $(u,v)$  is not in  $E$  assume  $c(u,v)=0$ .
  - We have source  $s$  and sink  $t$ .
  - Assume that every vertex  $v$  in  $V$  is on some path from  $s$  to  $t$ .
- Following is an illustration of a network flow:



$$c(s,v1)=16$$

$$c(v1,s)=0$$

$$c(v2,s)=0 \dots$$

# Conditions for Network Flow

For each edge  $(u,v)$  in  $E$ , the flow  $f(u,v)$  is a real valued function that must satisfy following 3 conditions :

- Capacity Constraint :  $\forall u,v \in V, f(u,v) \leq c(u,v)$
- Skew Symmetry :  $\forall u,v \in V, f(u,v) = -f(v,u)$
- Flow Conservation:  $\forall u \in V - \{s,t\} \sum_{v \in V} f(s,v) = 0$

Skew symmetry condition implies that  $f(u,u)=0$ .

## The Value of a Flow.

The value of a flow is given by :

$$|f| = \sum_{v \in V} f(s, v) = \sum_{v \in V} f(v, t)$$

The flow into the node is same as flow going out from the node and thus the flow is conserved. Also the total amount of flow from source  $s$  = total amount of flow into the sink  $t$ .

## Example of a flow

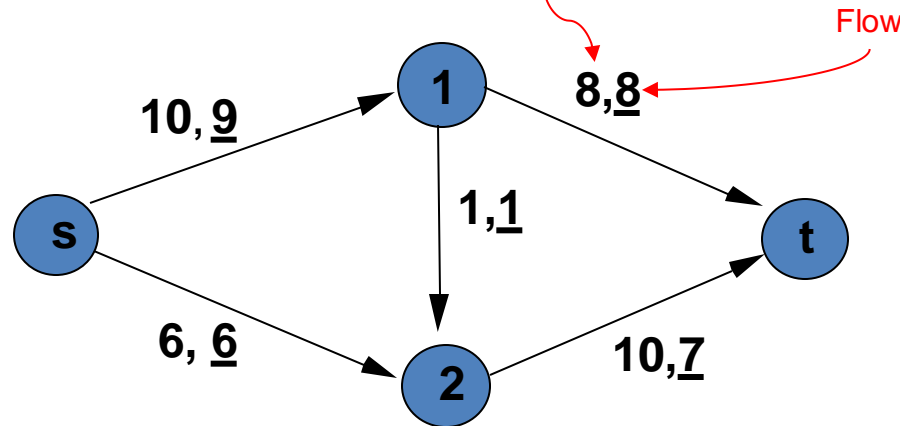


Table illustrating Flows and Capacity across different edges of graph above:

$f_{s,1} = 9$  ,  $c_{s,1} = 10$  (Valid flow since  $10 > 9$ )

$f_{s,2} = 6$  ,  $c_{s,2} = 6$  (Valid flow since  $6 \geq 6$ )

$f_{1,2} = 1$  ,  $c_{1,2} = 1$  (Valid flow since  $1 \geq 1$ )

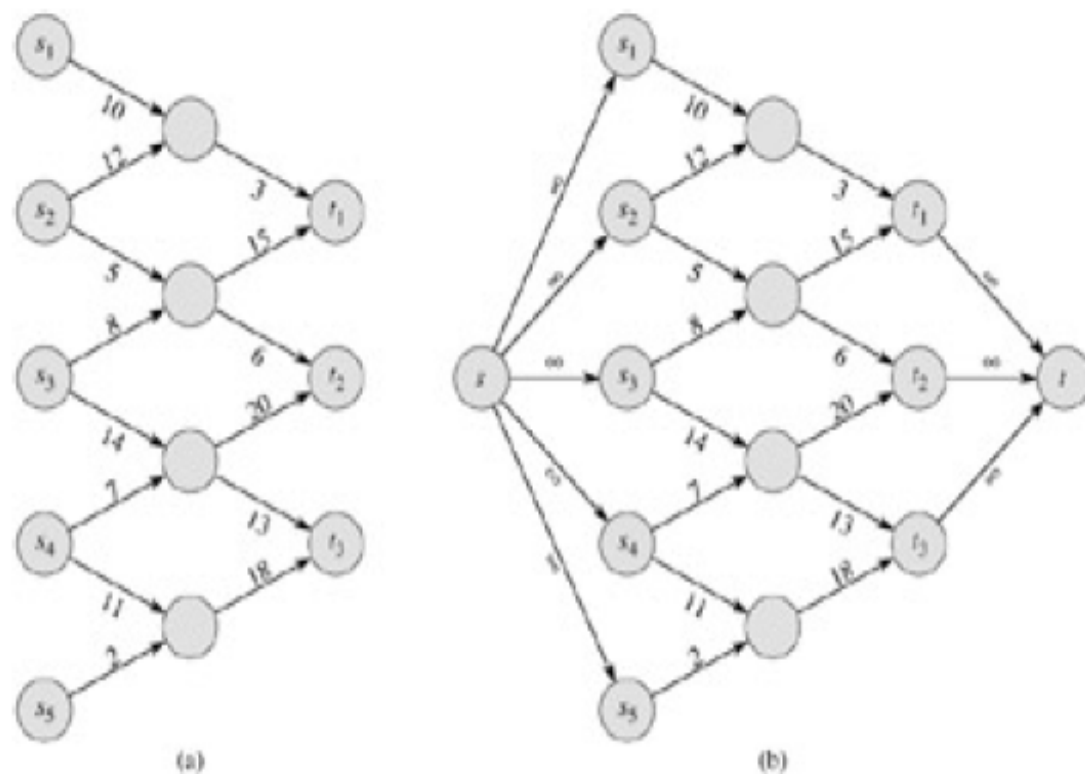
$f_{1,t} = 8$  ,  $c_{1,t} = 8$  (Valid flow since  $8 \geq 8$ )

$f_{2,t} = 7$  ,  $c_{2,t} = 10$  (Valid flow since  $10 > 7$ )

The flow across nodes 1 and 2 are also conserved as flow into them – flow out

# Networks with multiple sources and sinks

A maximum-flow problem may have several sources and sinks, rather than just one of each. The Lucky Company, for example, might actually have a set of  $m$  factories  $\{s_1, s_2, \dots, s_m\}$  and a set of  $n$  warehouses  $\{t_1, t_2, \dots, t_n\}$ .



Converting a multiple-source, multiple-sink maximum-flow problem into a problem with a single source and a single sink. (a) A flow network with five sources  $S = \{s_1, s_2, s_3, s_4, s_5\}$  and three sinks  $T = \{t_1, t_2, t_3\}$ . (b) An equivalent single-source, single-sink flow network. We add a supersource  $s$  and an edge with infinite capacity from  $s$  to each of the multiple sources. We also add a supersink  $t$  and an edge with infinite capacity from each of



We can reduce the problem of determining a maximum flow in a network with multiple sources and multiple sinks to an ordinary maximum-flow problem. [Figure 26.2\(b\)](#) shows how the network from (a) can be converted to an ordinary flow network with only a single source and a single sink. We add a *supersource*  $s$  and add a directed edge  $(s, s_i)$  with capacity  $c(s, s_i) = \infty$  for each  $i = 1, 2, \dots, m$ . We also create a new *supersink*  $t$  and add a directed edge  $(t_i, t)$  with capacity  $c(t_i, t) = \infty$  for each  $i = 1, 2, \dots, n$ . Intuitively, any flow in the network in (a) corresponds to a flow in the network in (b), and vice versa. The single source  $s$  simply provides as much flow as desired for the multiple sources  $s_i$ , and the single sink  $t$  likewise consumes as much flow as desired for the multiple sinks  $t_i$ . [Exercise 26.1-3](#) asks you to prove formally that the two problems are equivalent.

# The Ford Fulkerson Method

The Ford-Fulkerson method is iterative. We start with  $f(u, v) = 0$  for all  $u, v \in V$ , giving an initial flow of value 0. At each iteration, we increase the flow value by finding an "augmenting path," which we can think of simply as a path from the source  $s$  to the sink  $t$  along which we can send more flow, and then augmenting the flow along this path. We repeat this process until no augmenting path can be found. The max-flow min-cut theorem will show that upon termination, this process yields a maximum flow.

FORD-FULKERSON-METHOD( $G, s, t$ )

```
1 initialize flow  $f$  to 0
2 while there exists an augmenting path  $p$ 
3     do augment flow  $f$  along  $p$ 
4 return  $f$ 
```

# Serial Algorithms & parallel algorithms

- **Serial Algorithms:** Suitable for running on an uniprocessor computer in which only one instruction executes at a time.
- **Parallel Algorithms:** Run on a multiprocessor computer that permits multiple execution to execute concurrently.

# PARALLEL COMPUTERS

- Computers with multiple processing units.
- They can be:
  - **Chip Multiprocessors:** Inexpensive laptops/desktops. They contain a single multicore integrated-circuit that houses multiple processor “cores” each of which is a full-fledged processor with access to common memory.

# PARALLEL COMPUTERS

- Computers with multiple processing units.
- They can be:
  - **Clusters:** Build from individual computers with a dedicated network system interconnecting them. Intermediate price/performance.

# PARALLEL COMPUTERS

- Computers with multiple processing units.
- They can be:
  - **Supercomputers:** Combination of custom architectures and custom networks to deliver the highest performance (instructions per second).  
High price.

# Models for parallel computing

- Although the random-access machine model was early accepted for serial computing, no model has been established for parallel computing.
- A major reason is that vendors have not agreed on a single architectural model for parallel computers.

# Models for parallel computing

- For example some parallel computers feature **shared memory** where all processors can access any location of memory.
- Others employ **distributed memory** where each processor has a private memory.
- However, the trend appears to be toward **shared memory multiprocessor**.



# Static threading

- Shared-memory parallel computers use **static threading**.
- Software abstraction of “virtual processors” or threads sharing a common memory.
- Each thread can execute code independently.
- For most applications, threads persist for the duration of a computation.

# PROBLEMS OF STATIC THREADING

- Programming a shared-memory parallel computer directly using static threads is difficult and error prone.
- Dynamically partitioning the work among the threads so that each thread receives approximately the same load turns out to be complicated.

# PROBLEMS OF STATIC THREADING

- The programmer must use complex communication protocols to implement a scheduler to load-balance the work.
- This has led to the creation of **concurrency platforms**. They provide a layer of software that coordinates, schedules and manages the parallel-computing resources.

# DYNAMIC MULTITHREADING

- Class of concurrency platform.
- It allows programmers to specify parallelism in applications without worrying about communication protocols, load balancing, etc.
- The concurrency platform contains a scheduler that load-balances the computation automatically.

# DYNAMIC MULTITHREADING

- It supports:
  - **Nested parallelism:** It allows a subroutine to be spawned, allowing the caller to proceed while the spawned subroutine is computing its result.
  - **Parallel loops:** regular for loops except that the iterations can be executed concurrently.

# ADVANTAGES OF DYNAMIC MULTITHREADING

- The user only specifies the logical parallelism.
- Simple extension of the serial model with:  
**parallel, spawn** and **sync**.
- Clean way to quantify parallelism.
- Many multithreaded algorithms involving nested parallelism follow naturally from the Divide & Conquer paradigm.

# BASICS OF MULTITHREADING

- Fibonacci Example

- The serial algorithm: Fib( $n$ )
- Repeated work
- Complexity
- However, recursive calls are independent!
- Parallel algorithm: P-Fib( $n$ )

FIB( $n$ )

1    **if**  $n \leq 1$

2        **return**  $n$

3    **else**  $x = \text{FIB}(n - 1)$

4         $y = \text{FIB}(n - 2)$

5        **return**  $x + y$

# Serialization

- Concurrency keywords: **spawn**, **sync** and **parallel**
- The serialization of a multithreaded algorithm is the serial algorithm that results from deleting the concurrency keywords.



# NESTED PARALLELISM

- It occurs when the keyword **spawn** precedes a procedure call.
- It differs from the ordinary procedure call in that the procedure instance that executes the spawn - **the parent** – may continue to execute in parallel with the spawn subroutine – **its child** - instead of waiting for the child to complete.

# Keyword spawn

- It doesn't say that a procedure **must** execute concurrently with its **spawned** children; only that it **may**!
- The concurrency keywords express the **logical parallelism** of the computation.
- At runtime, it is up to the **scheduler** to determine which subcomputations actually run concurrently by assigning them to processors.

# Keyword sync

- A procedure cannot safely use the values returned by its spawned children until after it executes a **sync** statement.
- The keyword **sync** indicates that the procedure must wait until all its spawned children have been completed before proceeding to the statement after the sync.
- Every procedure executes a **sync** implicitly before it returns.

# FIB procedure to use dynamic multithreading

P-FIB( $n$ )

```
1  if  $n \leq 1$   
2      return  $n$   
3  else  $x = \text{spawn P-FIB}(n - 1)$   
4       $y = \text{P-FIB}(n - 2)$   
5      sync  
6      return  $x + y$ 
```

# Computational dag

- We can see a **multithread computation** as a directed acyclic graph  $G=(V,E)$  called a **computational dag**.
- The vertices are instructions and the edges represent dependencies between instructions, where  $(u,v) \in E$  means that instruction  $u$  must execute before instruction  $v$ .

# Computational dag

- If a chain of instructions contains no parallel control (no **spawn**, **sync**, or **return**), we may group them into a single ***strand***, each of which represents one or more instructions.
- Instructions involving parallel control are not included in strands, but are represented in the structure of the dag.

# Computational dag

- For example, if a strand has two successors, one of them must have been spawned, and a strand with multiple predecessors indicates the predecessors joined because of a **sync**.
- Thus, in the general case, the set  $V$  forms the set of strands, and the set  $E$  of directed edges represents dependencies between strands induced by parallel control.

# Computational dag

- If  $G$  has a directed path from strand  $u$  to strand  $v$ , we say that the two strands are *(logically) in series*. Otherwise, strands  $u$  and  $v$  are *(logically) in parallel*.
- We can picture a multithreaded computation as a dag of strands embedded in a tree of procedure instances.
- Example!



# Computational dag

- We can classify the edges:
  - ***Continuation edge*** : connects a strand  $u$  to its successor  $u'$  within the same procedure instance.
  - **Call edges**: representing normal procedure calls.
  - **Return edges**: When a strand  $u$  returns to its calling procedure and  $x$  is the strand immediately following the next **sync** in the calling procedure.
- A computation starts with an **initial strand** and ends with a single **final strand**.